

AASD 4004

Machine Learning - II

Applied AI Solutions Developer Program



Module 1

NLP Revisited

Vejey Gandyer

Agenda

What is NLP

NLP Tasks

Applications

NLP Building Blocks

Challenges of NLP

NLP Approaches

NLP Pipeline

NLP

What is it?



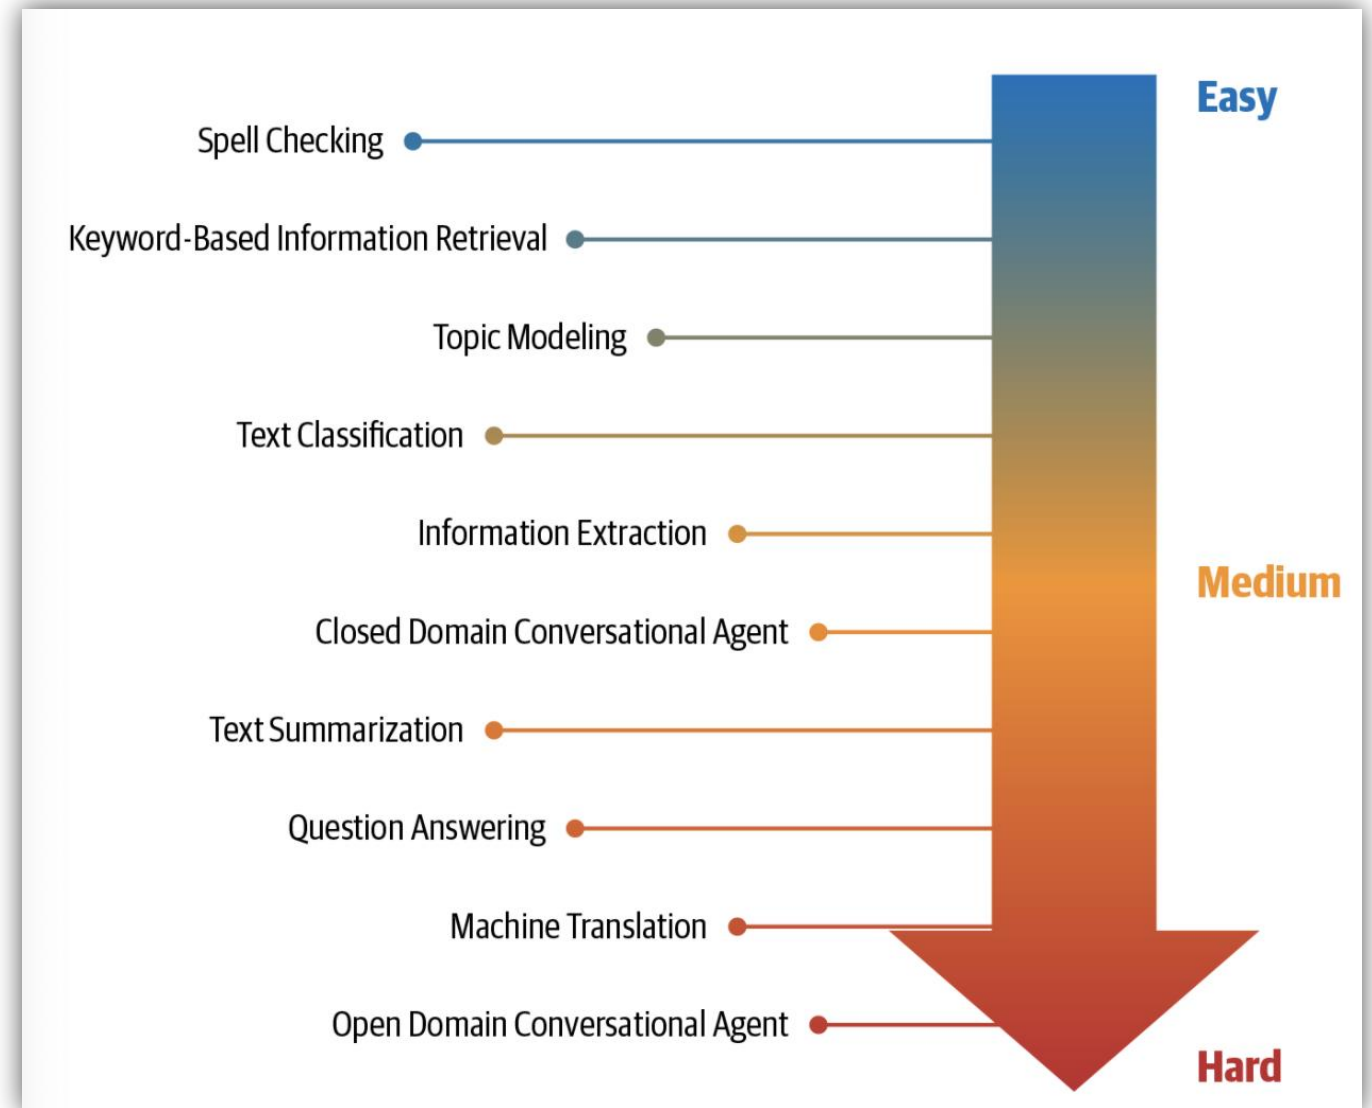
Natural Language Processing (NLP)

Natural language processing (NLP) is a subfield of linguistics, computer science, and artificial intelligence concerned with the interactions between computers and human language, how to program computers to process and analyze large amounts of natural language data.

NLP Tasks

NLP Tasks

Language Modeling
Text Classification
Information Extraction
Topic Modeling
Information Retrieval
Text Summarization
Question Answering
Machine Translation
Conversational Agent



NLP Applications

Text Classification

- Predict Tags or Categories
- Predict Sentiment
- Filter Spam mails


Sequence Applications

- Part Of Speech Tags
- Named Entity Recognition
- Semantic Slots

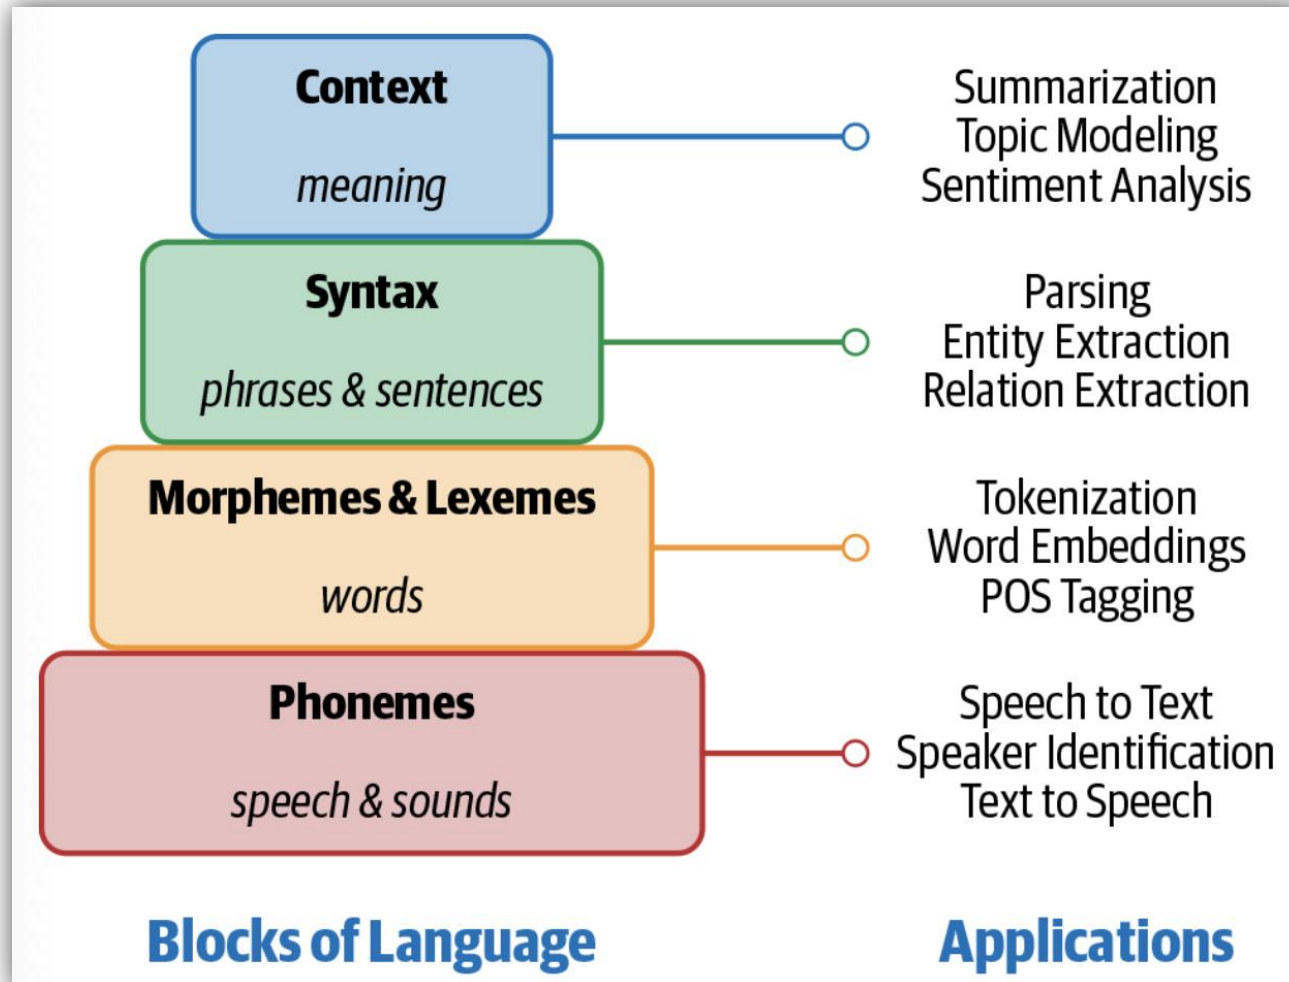
Sequence to Sequence

- Machine Translation
- Summarization
- Speech Recognition
- Question Answering

NLP Building Blocks

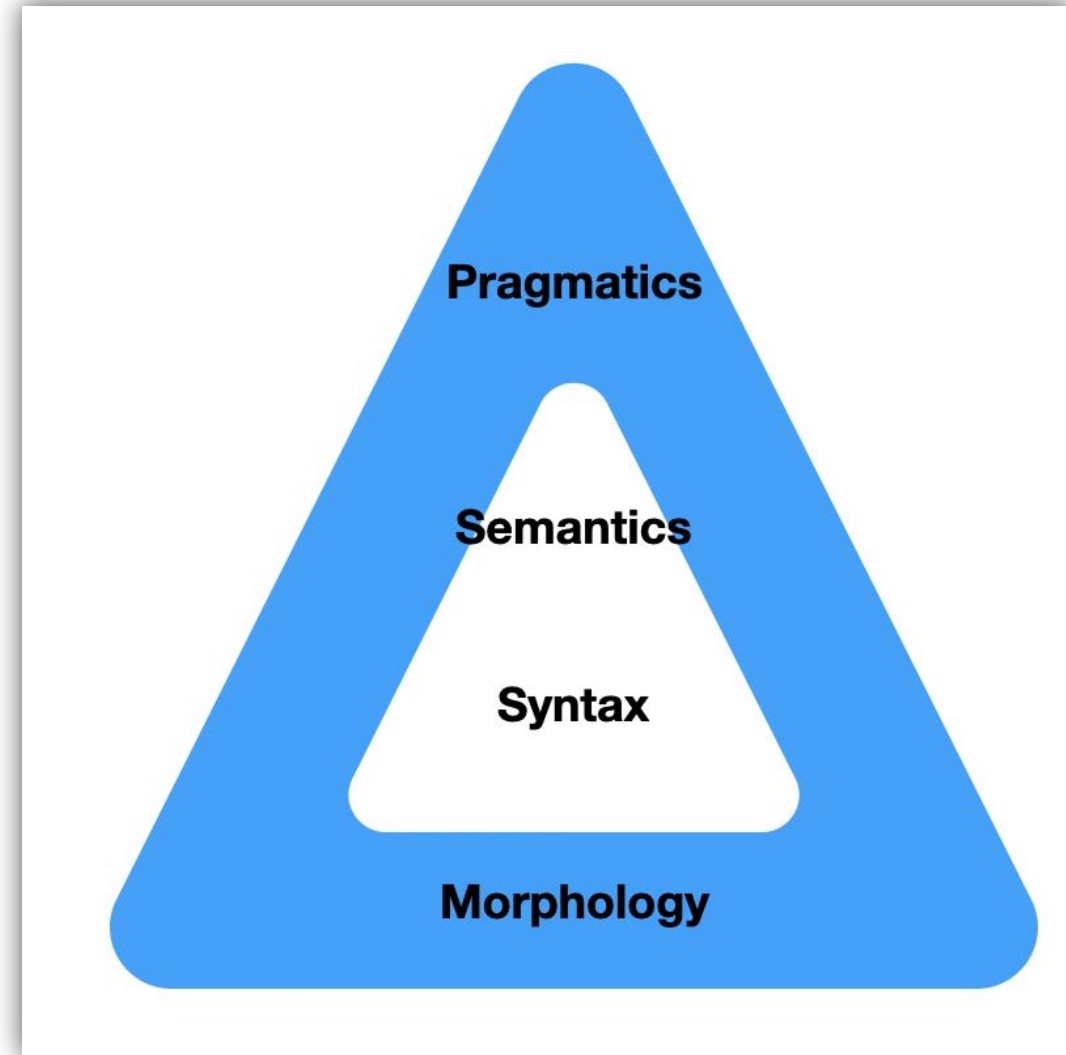


Building Blocks of Language



Linguistic Pyramid

- Morphology - Pre-processing
- Syntax
- Semantics
- Pragmatics



Representations

- Word Embeddings
- Sentence Embeddings
- Topic Models (Documents)
- Vector Space Models
- Similarity Graphs

NLP Approaches

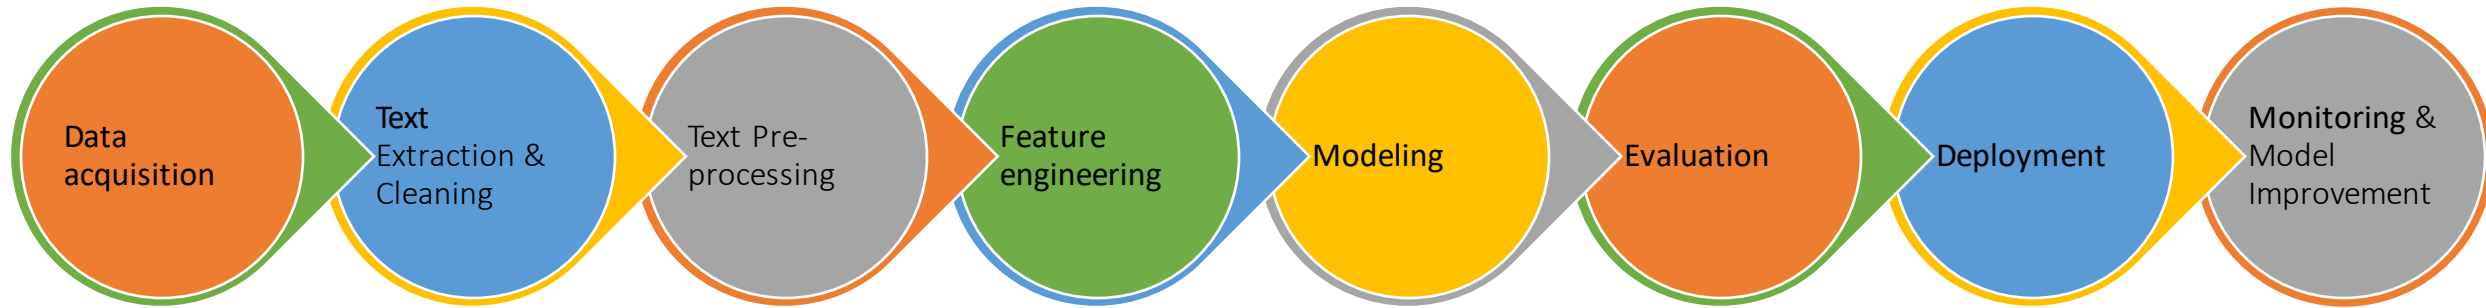
NLP Approaches

- Rule Based
 - Regular Expressions
 - Context-free Grammars
- Machine Learning
 - Probabilistic Modeling
 - Linear Classifiers
- Deep Learning
 - Recurrent Neural Networks
 - Convolutional Neural Networks

NLP Pipeline



NLP Pipeline



NLP Pipeline

Data Acquisition



Data Acquisition - Motivation

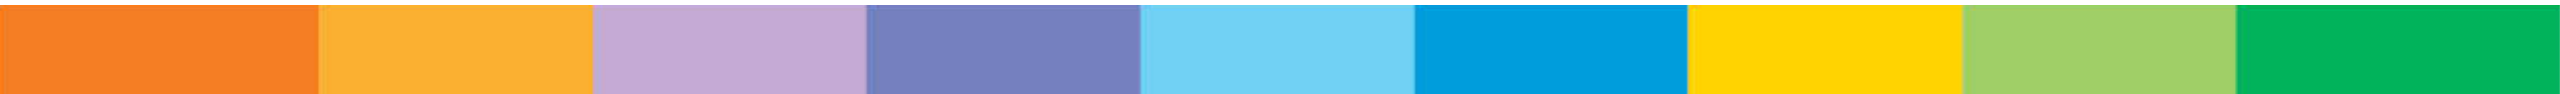
Data Acquisition – heart of any ML System

Hypothetical NLP project: Creating a NLP system to identify whether the user query is a sales query or a customer care query.

Ideal scenario: Millions of data points available already. No need for data acquisition

Real scenario: No data or less data with unlabelled classes. Need to acquire data

But how???



Data Acquisition

Use a Public Dataset

Scrape data from ungated websites

Product Intervention

Data Augmentation

- Synonym Replacement

- Back Translation

- TF-IDF Word Replacement

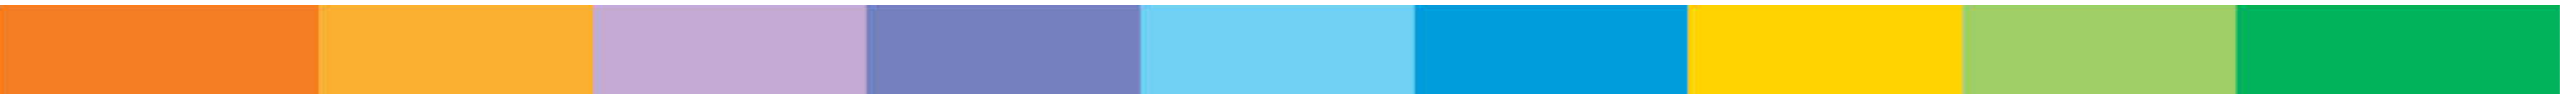
- Bigram Flipping

- Replacing entities

- Adding noise

NLP Pipeline

Text Extraction & Cleaning



Text Extraction and Cleaning - Motivation


Text Extraction - Process of extracting raw text from input data by removing all the other unwanted non-textual information (markup, metadata, ...)

Text Cleaning – a must as real-world data is always messy, noisy and incomplete 99% of the time

Ideal scenario: Cleaned data points already. No need for data cleaning

Real scenario: Noisy, incomplete, messy data. Need to clean data

But how???



Text Extraction and Cleaning

Various formats of input data (PDF, HTML, continuous stream, ...)

[illegible]

Supplement Facts

Serving Size 1 Lozenge

Amount Per Serving	% Daily Value
Vitamin B ₁₂ 5000 mcg (as methylcobalamin)	208333%

OTHER INGREDIENTS: XYLITOL, MANNITOL, STEARIC ACID, CITRIC ACID, NATURAL RASPBERRY FLAVOR, NATURAL BERRY FLAVOR, MAGNESIUM STEARATE.

DISTRIBUTED BY:
AMAZON.COM SERVICES, INC.
410 TERRY AVENUE N.,
SEATTLE, WA 98109

PRODUCT OF: CHINA

**MADE IN THE U.S.A. FROM U.S.
AND IMPORTED INGREDIENTS.**

CAUTION: AS WITH ANY DIETARY SUPPLEMENT, CONSULT YOUR HEALTHCARE PRACTITIONER BEFORE USING THIS PRODUCT, ESPECIALLY IF YOU ARE PREGNANT, NURSING, ANTICIPATE SURGERY, TAKE MEDICATION ON A REGULAR BASIS OR ARE OTHERWISE UNDER MEDICAL SUPERVISION.

**KEEP OUT OF REACH OF CHILDREN.
FOR ADULT USE ONLY.**

DO NOT USE IF SEAL IS BROKEN
OR MISSING.

STORE IN A COOL, DRY PLACE.

SUGGESTED USE:
TAKE 1 LOZENGE DAILY.

```
<a
href="https://www.oreilly.com/library/view/practical-natural-language/97
81492054047/">

</a>

<h1 style="color: #e74c3c;">Commonly Asked Questions</h1>
<ul>
<li> Can I contribute to the book?
<p>The book is accompanied by open source Jupyter notebooks and demo
applications. If you are a great ML or front-end engineer looking to
build something meaningful you can apply by filling <a
href="https://goo.gl/forms/do6NcW251iX26ajk1">this form</a>. Also refer
to the next question.
```

stackoverflow About Products For Teams Search...

Home PUBLIC Stack Overflow Tags Users FIND A JOB Jobs Companies TEAMS What's this? Free 30 Day Trial

How to get the current time in Python

Asked 11 years, 11 months ago Active 16 days ago Viewed 3.5m times

What is the module/method used to get the current time?

3116 python datetime time

share improve this question follow

edited Dec 26 '19 at 15:57 Georgy 5,079 5 39 49

asked Jan 6 '09 at 4:54 user46646 128k 43 73 82

15 please note, the most voted answers are for timezone-naïve datetime, while we see that in production environment more and more services across the world are connected together and timezone-aware datetime become the required standard – Slawomir Lenart Apr 29 at 17:12

also, for django, see [here](#) for timezone-aware – Anupam Sep 19 at 6:42

add a comment

43 Answers Active Oldest Votes

1 2 Next

Use:

3262

```
>>> import datetime
>>> datetime.datetime.now()
datetime.datetime(2009, 1, 6, 15, 8, 24, 78915)

>>> print(datetime.datetime.now())
2009-01-06 15:08:24.789150
```

And just the time:

```
>>> datetime.datetime.now().time()
datetime.time(15, 8, 24, 78915)

>>> print(datetime.datetime.now().time())
15:08:24.789150
```

See [the documentation](#) for more information.

To save typing, you can import the `datetime` object from the `datetime` module:

```
>>> from datetime import datetime
```

Then remove the leading `datetime.` from all of the above.

Web scraping

Text =
What is the module/method used to get the current time?

Answer =
Use:

```
>>> import datetime
>>> datetime.datetime.now()
datetime.datetime(2009, 1, 6, 15, 8, 24, 78915)
```

```
>>> print(datetime.datetime.now())
2009-01-06 15:08:24.789150
```

And just the time:

```
>>> datetime.datetime.now().time()
datetime.time(15, 8, 24, 78915)
```

```
>>> print(datetime.datetime.now().time())
15:08:24.789150
```

See the documentation for more information.
To save typing, you can import the `datetime` object from the `datetime` module:

```
>>> from datetime import datetime
```

Then remove the leading `datetime.` from all of the above.

Text extraction from Scanned images

In the nineteenth century the only kind of linguistics considered seriously was this comparative and historical study of words in languages known or believed to be *cognate*—say the Semitic languages, or the Indo-European languages. It is significant that the Germans who really made the subject what it was, used the term *Indo-germanisch*. Those who know the popular works of Otto Jespersen will remember how firmly he declares that linguistic science is historical. And those who have noticed

PyPDF, PDFMiner, ...
can be used to extract
PDF text

If PDF text is a scanned
image, use **Tesseract**
library

```
from PIL import Image
from pytesseract import image_to_string
filename = "somefile.png"
text = image_to_string(Image.open(filename))
print(text)
```

```
'in the nineteenth century the only Kind of linguistics considered\nseriously
was this comparative and historical study of words in languages\nknown or
believed to Fe cognate—say the Semitic languages, or the Indo-\nEuropean
languages. It is significant that the Germans who really made\nthe subject what
it was, used the term Indo-germanisch. Those who know\nthe popular works of
Otto Jespersen will remember how fitmly he\ndeclares that linguistic
science is historical. And those who have noticed'
```

Unicode removal

↑ U+2191	🙄 U+1F647	— U+2010	\, U+FF64	⊖ U+0398	♥ U+1F49A	😊 U+263B	♂ U+056E	ℷ U+0C2C	୪ U+0CA0
γ U+03B3	ᑭ U+12CE	♥ U+1F49C	μ U+03BC	🚀 U+1F680	🎵 U+266A	☾ U+FE36	· U+30FB	♡ U+10E6	” U+2036
⚙ U+263C	। U+0964	○ U+26AC	ह U+0939	且 U+4E14	ब U+09F0	₵ U+0F4F	‰ U+2030	テ U+30C7	↩ U+21A9
‘ U+2018	” U+2033	ℙ U+026A	୪ U+0DA2	😂 U+1F639	Δ U+0394	ù U+00F9	↩ U+27AB	ÿ U+0177	🧘 U+1F9D8

To remove non-textual symbols and special characters, use Unicode Normalization

Use `string.encode("utf-8")`

```
text = 'I love 🍕! Shall we book a 🚗 to gizza?'
Text = text.encode("utf-8")
print(Text)
```

`b'I love Pizza \xf0\x9f\x8d\x95! Shall we book a cab \xf0\x9f\x9a\x95 to get pizza?'`

Spelling correction

```
import requests
import json

api_key = "<ENTER-KEY-HERE>"
example_text = "Hollo, wrld" # the text to be spell-checked

data = {'text': example_text}
params = {
    'mkt': 'en-us',
    'mode': 'proof'
}
headers = {
    'Content-Type': 'application/x-www-form-urlencoded',
    'Ocp-Apim-Subscription-Key': api_key,
}

response = requests.post(endpoint, headers=headers, params=params, data=data)
json_response = response.json()
print(json.dumps(json_response, indent=4))
```

Microsoft has APIs for spell checking

Use **Levenshtein distance** for detecting and correcting spelling mistakes

```
"suggestions": [
  {
    "suggestion": "Hello",
    "score": 0.9115257530801
  },
  {
    "suggestion": "Hollow",
    "score": 0.858039839213461
  },
  {
    "suggestion": "Hallo",
    "score": 0.597385084464481
  }
]
```

NLP Pipeline

Text Pre-processing



Text Pre-processing - Motivation

Text Pre-processing - Process of preparing raw text extracted from data sources by some processes like Sentence Segmentation, Word Tokenization, Stop words, Stemming & Lemmatization, Special characters removal, POS tagging, Coreference resolution, etc.

Text Processing – a must as real-world needs to be in a certain format for a machine learning algorithm to be accepted as input

But how???



Text Pre-processing

Sentence Segmentation

Word Tokenization

Stop words removal

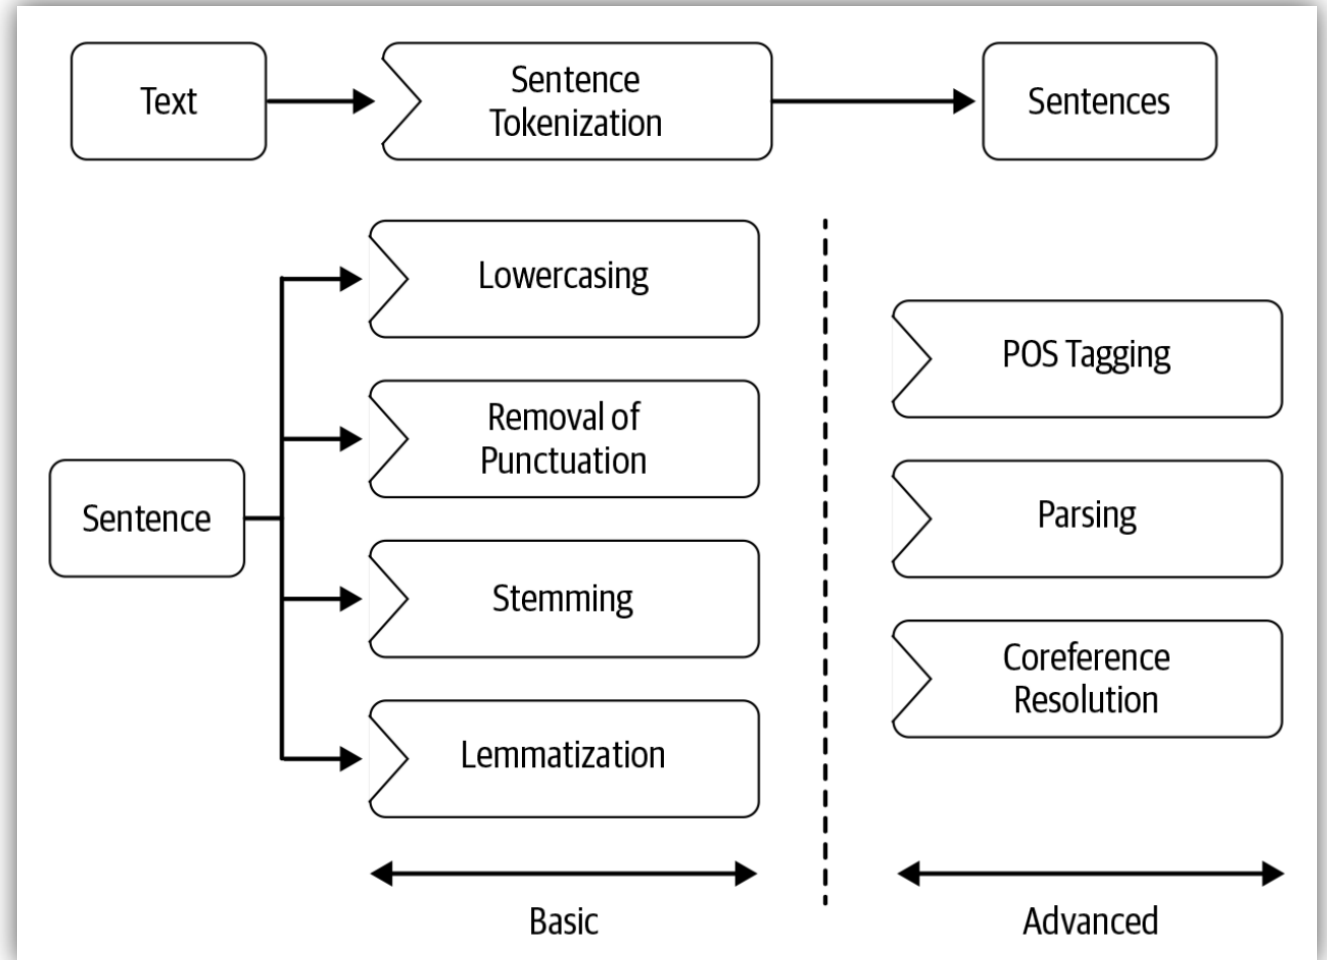
Stemming & Lemmatization

Normalization

POS tagging

Parse Tree

Coreference resolution



Sentence Segmentation

```
mytext = """In the previous chapter, we saw examples of some common NLP applications that we might encounter in everyday life. If we were asked to build such an application, think about how we would approach doing so at our organization. We would normally walk through the requirements and break the problem down into several sub-problems, then try to develop a step-by-step procedure to solve them. Since language processing is involved, we would also list all the forms of text processing needed at each step. This step-by-step processing of text is known as pipeline. """
```

Breaking a big document text into sentences

```
from nltk.tokenize import sent_tokenize
my_sentences = sent_tokenize(mytext)
print(my_sentences)
```

```
['In the previous chapter, we saw examples of some common NLP applications that we might encounter in everyday life.',  
'If we were asked to build such an application, think about how we would approach doing so at our organization.',  
'We would normally walk through the requirements and break the problem down into several sub-problems, then try to develop  
'Since language processing is involved, we would also list all the forms of text processing needed at each step.',  
'This step-by-step processing of text is known as pipeline.'  
]
```

Word Tokenization

```
['In the previous chapter, we saw examples of some common NLP applications that we might encounter in everyday life.',  
'If we were asked to build such an application, think about how we would approach doing so at our organization.',  
'We would normally walk through the requirements and break the problem down into several sub-problems, then try to develop  
'Since language processing is involved, we would also list all the forms of text processing needed at each step.',  
'This step-by-step processing of text is known as pipeline.'  
]
```

```
from nltk.tokenize import word_tokenize  
for sentence in my_sentences:  
    print(sentence)  
    print(word_tokenize(sentence))
```

Breaking a big
sentence into
words or
tokens

This step-by-step processing of text is known as pipeline.

```
['This', 'step-by-step', 'processing', 'of', 'text', 'is', 'known', 'as', 'pipeline', '.']
```

Stop words removal

"In the previous chapter, we saw examples of some common NLP applications that we might encounter in everyday life."

```
from nltk.corpus import stopwords
from string import punctuation
def preprocess_corpus(texts):
    mystopwords = set(stopwords.words("english"))
    def remove_stops_digits(tokens):
        return [token.lower() for token in tokens if token not in mystopwords and
                not token.isdigit() and token not in punctuation]
    return [remove_stops_digits(word_tokenize(text)) for text in texts]
```

['in', 'previous', 'chapter', 'saw', 'examples', 'common', 'nlp', 'applications', 'might', 'encounter', 'everyday', 'life']

Stemming

Removes suffixes
and reduces a word
to some **base form**

cars revolution

```
from nltk.stem.porter import PorterStemmer
stemmer = PorterStemmer()
word1, word2 = "cars", "revolution"
print(stemmer.stem(word1), stemmer.stem(word2))
```

car revolut

Lemmatization

better

```
from nltk.stem import WordNetLemmatizer
lemmatizer = WordNetLemmatizer()
print(lemmatizer.lemmatize("better", pos="a"))
```

good

Also removes suffixes and reduces a word to some **base form** or **lemma**

Its lemma will be there in the dictionary (meaningful)

Stemming

adjustable -> adjust
formality -> formaliti
formaliti -> formal
airliner -> airlin

Lemmatization

was -> (to) be
better -> good
meeting -> meeting

Advanced pre-processing

To use spacy, first install them in your
conda / miniconda / virtual environment

pip3 install spacy

*python3 -m spacy download
en_core_web_sm*

```
import spacy
nlp = spacy.load('en_core_web_sm')
doc = nlp(u"""Charles Spencer Chaplin was born on 16 April 1889
to Hannah Chaplin (born Hannah Harriet Pedlingham Hill)
and Charles Chaplin Sr""")
for token in doc:
    print(token.text, token.lemma_, token.pos_,
          token.shape_, token.is_alpha, token.is_stop)
```

Input

Chaplin wrote, directed, and composed the music for most of his films.

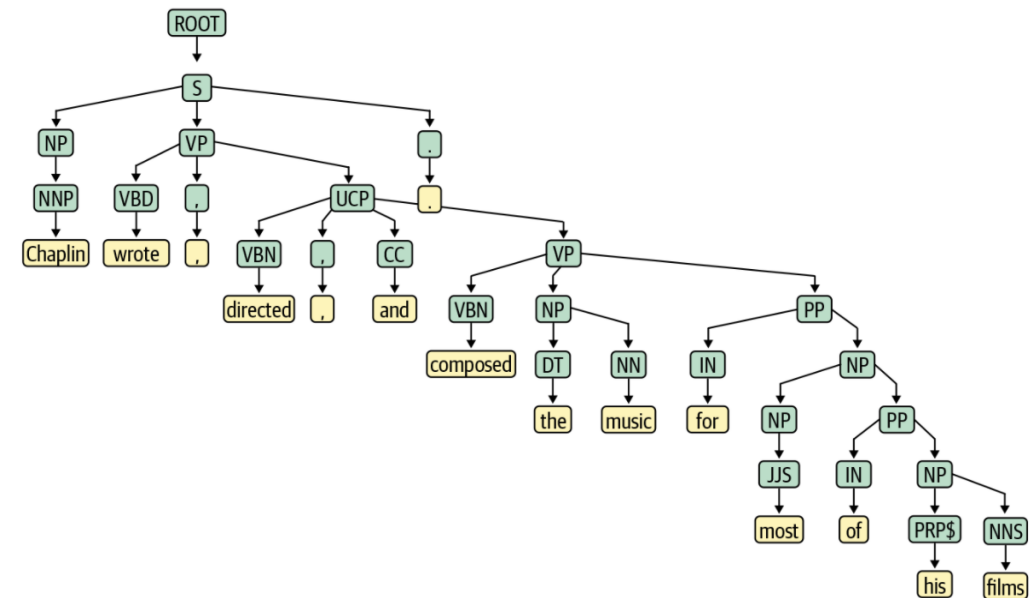
Tokenization with Lemmatization

Chaplin wrote, directed, and composed the music for most of his films.

POS Tagging

Chaplin wrote, directed, and composed the music for most of his films.

Parse Tree

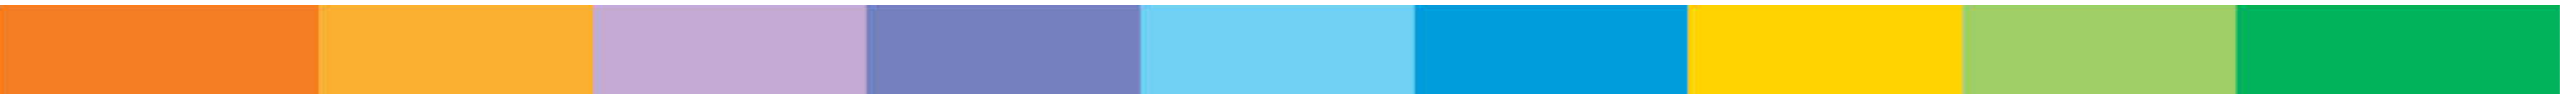


Coreference Resolution

Chaplin wrote, directed, and composed the music for most of his films.

NLP Pipeline

Feature Engineering



Feature Engineering - Motivation

Feature Engineering - Process of set of methods that will accomplish the task of feature extraction (converting text into numeric vectors)

Feature Engineering is dealt in the future lecture in detail.

Two major categories are defined here.

1. Classical NLP / ML Pipeline
2. DL Pipeline

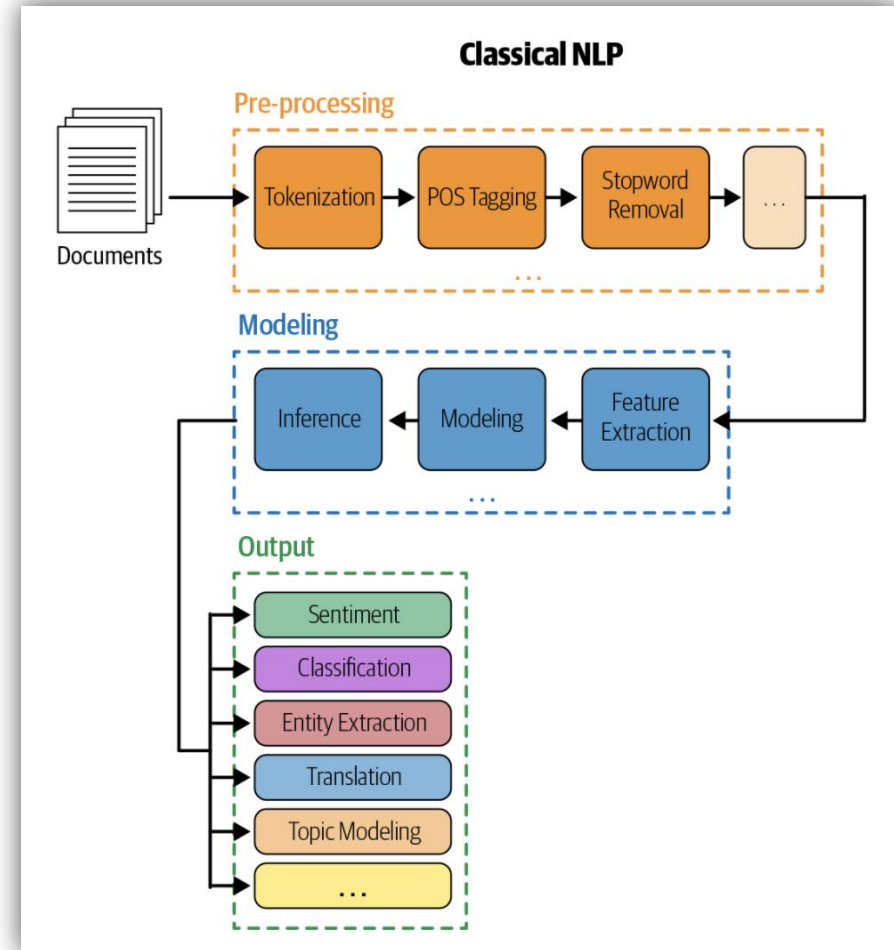
Feature Engineering - Classical NLP

Converts the raw data into a format that can be consumed by a machine.

In Classical ML, we have turned categorical variables into numbers and fed into the model.

In NLP, we need to convert the text into some form of numeric vectors and will feed it into the model

In Classical NLP, the feature extraction process is **handcrafted** and done by engineers who have **domain expertise** in the area of the problem in hand

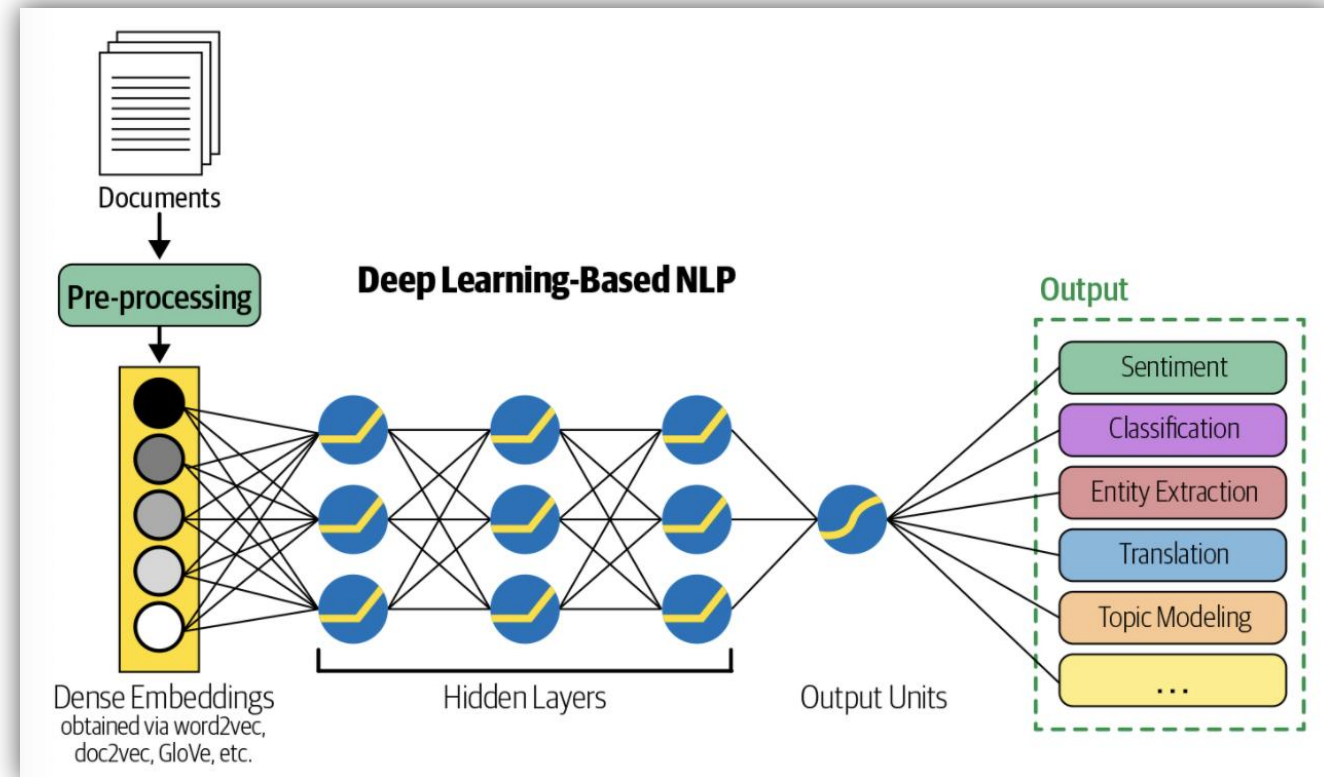


Feature Engineering - Deep NLP

Converts the raw data into a format that can be consumed by a machine.

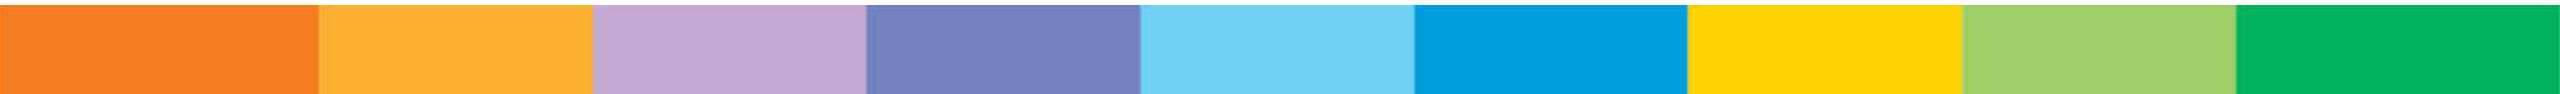
In Deep Learning NLP, the model takes care of the feature extraction process

Raw data after pre-processing is sent to the model directly



NLP Pipeline

Modeling



Modeling - Motivation

Modeling - Process of building a model depending on the amount of data we have in hand

Simple Heuristics - Regular Expressions, Rule-based approaches

Create a feature from the heuristic for the ML model

Pre-process input to the ML model



Modeling

Ensembles of Models

Feature Engineering

Transfer Learning

Data attribute	Decision path	Examples
Large data volume	<p>Can use techniques that require more data, like DL. Can use a richer set of features as well.</p> <p>If the data is sufficiently large but unlabeled, we can also apply unsupervised techniques.</p>	If we have a lot of reviews and metadata associated with them, we can build a sentiment-analysis tool from scratch.
Small data volume	<p>Need to start with rule-based or traditional ML solutions that are less data hungry. Can also adapt cloud APIs and generate more data with weak supervision.</p> <p>We can also use transfer learning if there's a similar task that has large data.</p>	This often happens at the start of a completely new project.
Data quality is poor and the data is heterogeneous in nature	More data cleaning and pre-processing might be required.	This entails issues like code mixing (different languages being mixed in the same sentence), unconventional language, transliteration, or noise (like social media text).
Data quality is good	Can directly apply off-the-shelf algorithms or cloud APIs more easily.	Legal text or newspapers.
Data consists of full-length documents	Choose the right strategy for breaking the document into lower levels, like paragraphs, sentences, or phrases, depending on the problem.	Document classification, review analysis, etc.

NLP Pipeline

Evaluation



Evaluation - Motivation

Evaluation – Measuring how good the model is

- 1) Use the right metric
- 2) Follow the right evaluation process

Intrinsic Evaluation

Extrinsic Evaluation



Intrinsic Evaluation

Metric	Description	Applications
Accuracy [48]	Used when the output variable is categorical or discrete. It denotes the fraction of times the model makes correct predictions as compared to the total predictions it makes.	Mainly used in classification tasks, such as sentiment classification (multiclass), natural language inference (binary), paraphrase detection (binary), etc.
Precision [48]	Shows how precise or exact the model's predictions are, i.e., given all the positive (the class we care about) cases, how many can the model classify correctly?	Used in various classification tasks, especially in cases where mistakes in a positive class are more costly than mistakes in a negative class, e.g., disease predictions in healthcare.
Recall [48]	Recall is complementary to precision. It captures how well the model can recall positive class, i.e., given all the positive predictions it makes, how many of them are indeed positive?	Used in classification tasks, especially where retrieving positive results is more important, e.g., e-commerce search and other information-retrieval tasks.
F1 score [49]	Combines precision and recall to give a single metric, which also captures the trade-off between precision and recall, i.e., completeness and exactness. F1 is defined as $(2 \times \text{Precision} \times \text{Recall}) / (\text{Precision} + \text{Recall})$.	Used simultaneously with accuracy in most of the classification tasks. It is also used in sequence-labeling tasks, such as entity extraction, retrieval-based questions answering, etc.

AUC [48]	Captures the count of positive predictions that are correct versus the count of positive predictions that are incorrect as we vary the threshold for prediction.	Used to measure the quality of a model independent of the prediction threshold. It is used to find the optimal prediction threshold for a classification task.
MRR (mean reciprocal rank) [50]	Used to evaluate the responses retrieved given their probability of correctness. It is the mean of the reciprocal of the ranks of the retrieved results.	Used heavily in all information-retrieval tasks, including article search, e-commerce search, etc.
MAP (mean average precision) [51]	Used in ranked retrieval results, like MRR. It calculates the mean precision across each retrieved result.	Used in information-retrieval tasks.
RMSE (root mean squared error) [48]	Captures a model's performance in a real-value prediction task. Calculates the square root of the mean of the squared errors for each data point.	Used in conjunction with MAPE in the case of regression problems, from temperature prediction to stock market price prediction.
MAPE (mean absolute percentage error) [52]	Used when the output variable is a continuous variable. It is the average of absolute percentage error for each data point.	Used to test the performance of a regression model. It is often used in conjunction with RMSE.

Intrinsic Evaluation

BLEU (bilingual evaluation understudy) [53]	Captures the amount of n-gram overlap between the output sentence and the reference ground truth sentence. It has many variants.	Mainly used in machine-translation tasks. Recently adapted to other text-generation tasks, such as paraphrase generation and text summarization.
METEOR [54]	A precision-based metric to measure the quality of text generated. It fixes some of the drawbacks of BLEU, such as exact word matching while calculating precision. METEOR allows synonyms and stemmed words to be matched with the reference word.	Mainly used in machine translation.
ROUGE [55]	Another metric to compare quality of generated text with respect to a reference text. As opposed to BLEU, it measures recall.	Since it measures recall, it's mainly used for summarization tasks where it's important to evaluate how many words a model can recall.
Perplexity [56]	A probabilistic measure that captures how confused an NLP model is. It's derived from the cross-entropy in a next word prediction task. The exact definition can be found at [56] .	Used to evaluate language models. It can also be used in language-generation tasks, such as dialog generation.

Extrinsic Evaluation

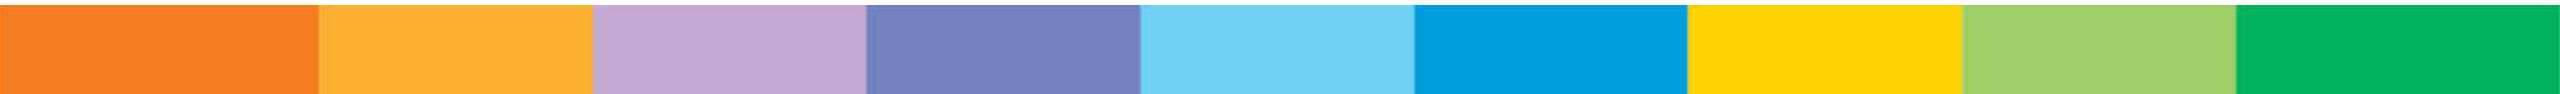
Involves the business metrics outside the AI/ML team

First, check to see if you achieve good intrinsic evaluation metrics

Then, go for extrinsic evaluation

NLP Pipeline

Deployment



Deployment

NLP model is deployed as a web service*

Some Cloud providers

Google Cloud Platform (GCP)

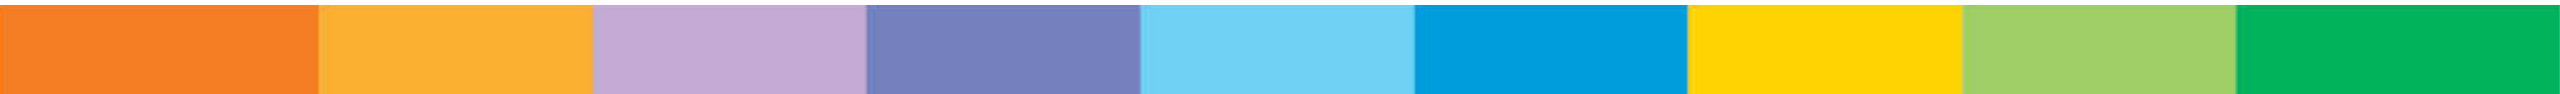
Amazon Web Services (AWS)

Microsoft Azure

* Will be seen in detail in Full Stack Data Science Systems course

NLP Pipeline

Monitoring & Model Improvement



Monitoring & Model Updation

Monitoring must be done on a constant real-time basis

Performance dashboards to be included in the project

Project attribute	Decision paths	Examples
More training data is generated post-deployment.	Once deployed, extracted signals can be used to automatically improve the model. Can also try online learning to train the model automatically on a daily basis.	Abuse-detection systems where users flag data.
Training data is not generated post-deployment.	Manual labeling could be done to improve evaluation and the models. Ideally, each new model has to be manually built and evaluated.	A subset of a larger NLP pipeline with no direct feedback.
Low model latency is required, or model has to be online with near-real-time response.	Need to use models that can be inferred quickly. Another option is to create memoization strategies like caching or have substantially bigger computing power.	Systems that need to respond right away, like any chatbot or an emergency tracking system.
Low model latency is not required, or model can be run in an offline fashion.	Can use more advanced and slower models. This can also help in optimizing costs where feasible.	Systems that can be run on a batch process, like retail product catalog analysis.

Further Reading

Practical Natural Language Processing

Soumya Vajjala, Anuj Gupta, Harshit Surana, Bodhisattwa Majumder