

# AASD 4004

# Machine Learning - II

Applied AI Solutions Developer Program



# Module 14

# Introduction to Audio Processing

Vejeý Gandyer



# Agenda

Audio Processing

Why

Basic Concepts

Tools for Audio Processing

Pyaudio

Librosa

Pydub

Applications

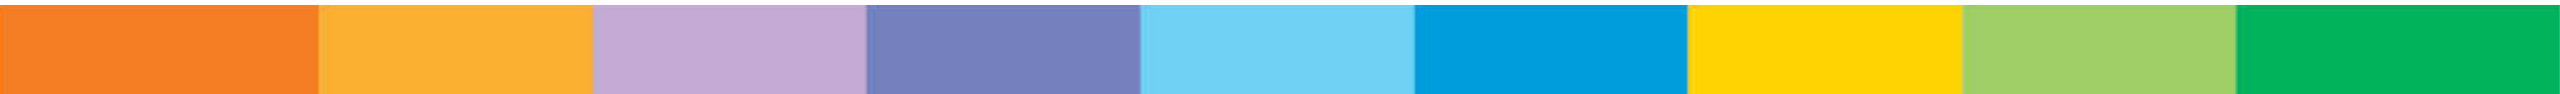
# Audio Processing

What is it?



# Audio Processing

Audio Processing means changing the characteristics of an audio signal in some way



# Why

# Why to process Audio?

To enhance audio

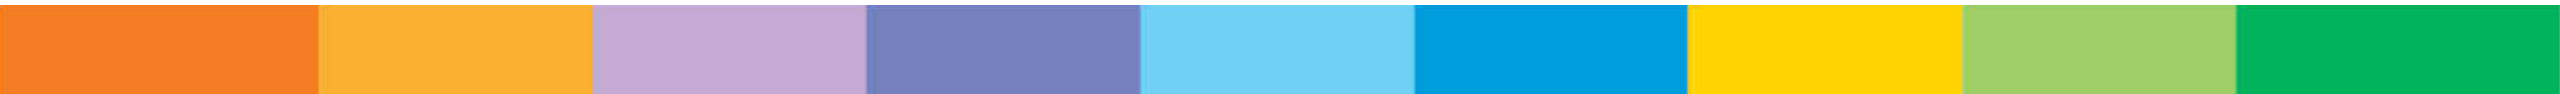
To separate channels

To create new sounds

To store sounds



# Basic Concepts





# Basic Concepts

Frequency


Sampling Rate

Sine wave

Amplitude

Discrete Fourier Transform - Calculates which frequencies are present, Converts time domain signal to a frequency domain

Noise Filtering




# Sound

Sound is a travelling vibration

2 basic attributes

Amplitude (loudness)

Frequency (measure of wave's vibration per unit time)




# Analog to Digital Conversion

Sampling - Convert the time-varying continuous signal  $x(t)$  to a discrete sequence  $x(n)$  of real numbers

Sampling period  $T_s$  - Interval between successive discrete samples

Sampling frequency -  $f_s = 1 / T_s$

Quantization - Replacing each real number of the sequence of samples with an approximation from a finite set of discrete values



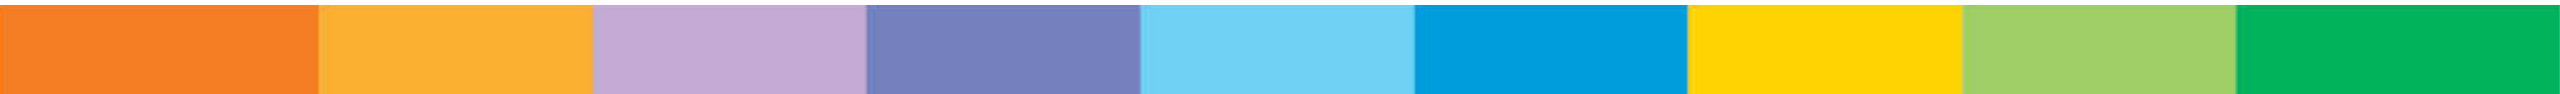
# Sine wave

## Sine wave formula

$$y(t) = A * \sin(2 * \pi * f * t)$$

# Frequency

Number of times a sine wave repeats a second

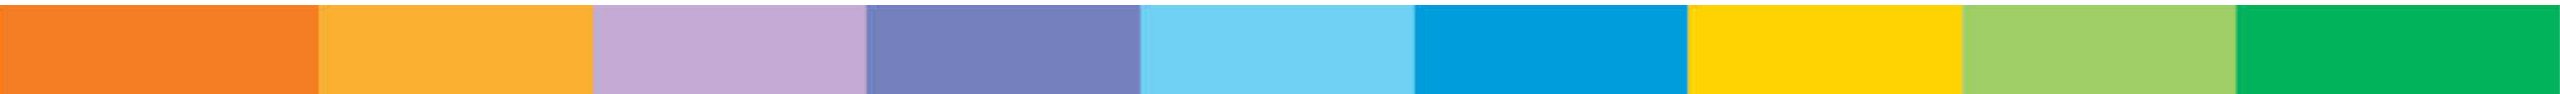


# Sampling rate

Real world signals are analog signals

In digitizing it, converting analog to digital, takes a sample at defined intervals

Eg: 48000



# Amplitude

Height of the sine wave

$$y(t) = A * \sin(2 * \pi * f * t)$$

# Let's create a sine wave

Create a sine wave

Write the sine wave in an audio file (.wav)

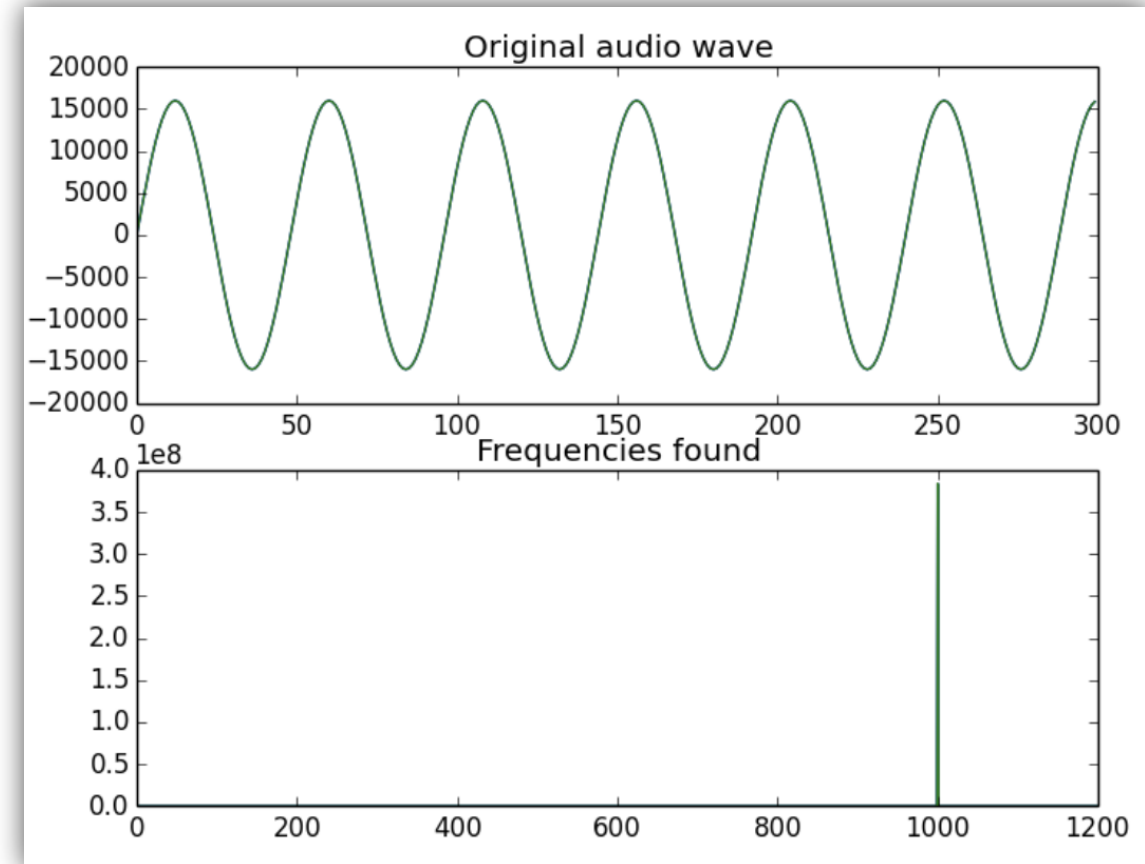
```
sine_wave = [np.sin(2 * np.pi * frequency * x/sampling_rate) for x in range(num_samples)]  
wav_file=wave.open(file, 'w')  
wav_file.setparams((nchannels, sampwidth, int(sampling_rate), nframes, comptype, compname))  
for s in sine_wave:  
    wav_file.writeframes(struct.pack('h', int(s*amplitude)))
```



# Discrete Fourier Transform (DFT)



Identifies what are the frequencies present in the signal

```
data_fft = np.fft.fft(data)
```



# Discrete Fourier Transform (DFT)

Identifies what are the frequencies present in the signal





# Cleaning a noisy signal (sine wave)

Generate a sine wave (signal)

Add noise

Filter the noise



# Cleaning a noisy signal (sine wave)


Generate a sine wave (signal)

Add noise

Filter the noise

```

frequency = 1000
noisy_freq = 50
num_samples = 48000
sampling_rate = 48000.0
    
```

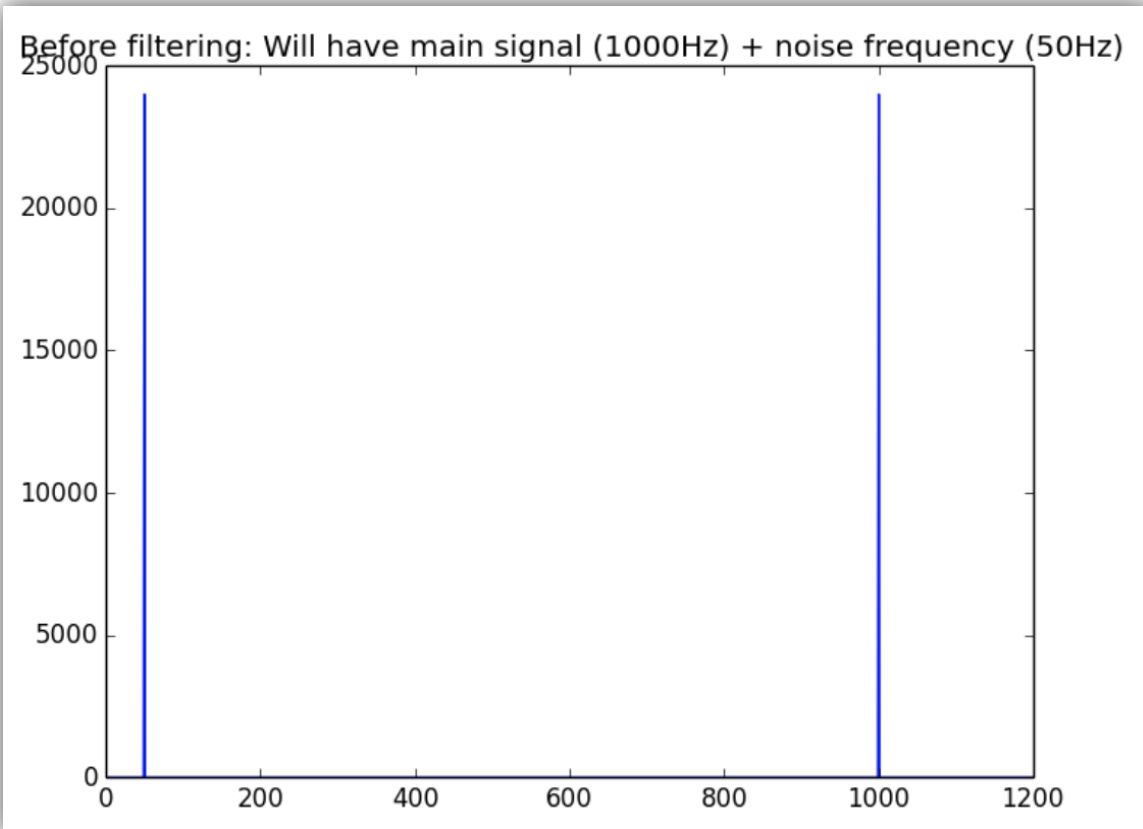


```

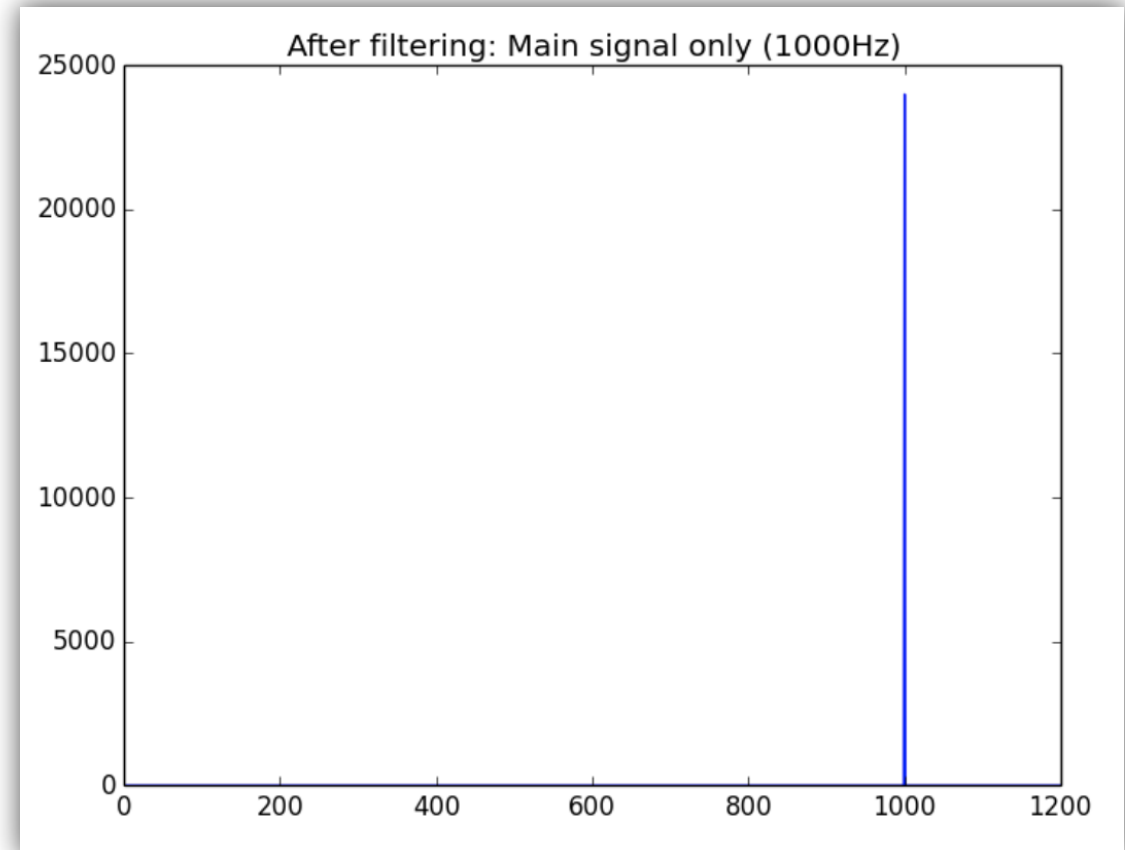
sine_wave = [np.sin(2 * np.pi * frequency * x1 / sampling_rate) for x1 in range(num_samples)]
sine_noise = [np.sin(2 * np.pi * noisy_freq * x1 / sampling_rate) for x1 in range(num_samples)]
    
```

# Cleaning a noisy signal (sine wave)

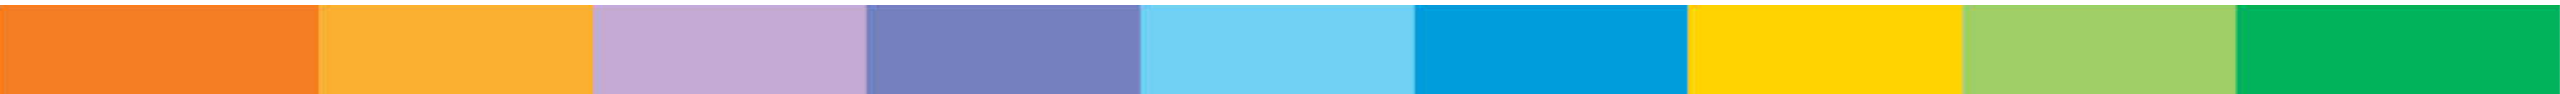
Noisy signal



Clean signal



# Tools for Audio Processing



# Tools

**ffmpeg/libav** - Handling multimedia files and streams

**sox** - Swiss Army knife of sound processing programs

**audacity** - Editing and playback software

**pyAudioAnalysis** - IO, advanced feature extraction and signal analysis

**librosa** - Nice library for audio analysis

**pydub** - Library for audio processing

# Pydub

<https://github.com/jiaaro/pydub>



# Loading a wav file

## pydub

```
from pydub import AudioSegment
import numpy as np
audiofile = AudioSegment.from_file("data/music_8k.mp3")
data_mp3 = np.array(audiofile.get_array_of_samples())
```

## scipy

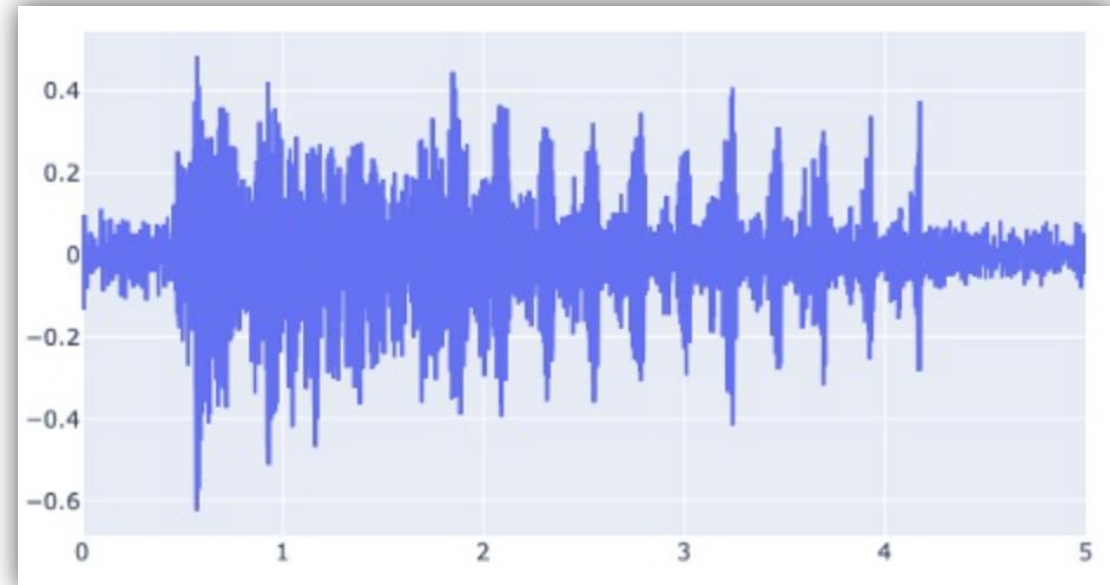
```
from scipy.io import wavfile
fs_wav, data_wav = wavfile.read("data/music_8k.wav")
```

# Normalization

```
fs_wav, data_wav = wavfile.read("data/lost_highway_small.wav")
data_wav_norm = data_wav / (2**15)
time_wav = np.arange(0, len(data_wav)) / fs_wav
plotly.offline.iplot({ "data": [go.Scatter(x=time_wav,
                                           y=data_wav_norm,
                                           name='normalized audio signal')]}))
```

Makes the signal values independent to the sample resolution

Squish the values of an audio signal in the  $(-1, 1)$  by dividing by  $2^{15}$




# Trim/Clip/Segment audio clips

```
data_wav_norm_crop = data_wav_norm[2 * fs_wav: 4 * fs_wav]
time_wav_crop = np.arange(0, len(data_wav)) / fs_wav
plotly.offline.iplot({ "data": [go.Scatter(x=time_wav_crop,
                                             y=data_wav_norm_crop,
                                             name='cropped audio signal')]}))
```

Clips the audio file portions by referring to the respective indices in the numpy array

Seconds need to be multiplied by the sampling frequency



[illegible]

# Remove silent segments from audio

Computes energy as the sum of squares of the samples


Calculates a threshold as 50% of the median energy value

Keeps the segments whose energy are above the set threshold

```
energies = [(s**2).sum() / len(s) for s in segments]
thres = 0.5 * np.median(energies)
index_of_segments_to_keep = (np.where(energies > thres)[0])
segments2 = segments[index_of_segments_to_keep]
new_signal = np.concatenate(segments2)
wavfile.write("data/obama_processed.wav", fs, new_signal)
```

# Remove silent segments from audio

Keeps the segments whose energy are above the set threshold




# Remove silent segments from audio

Original



Processed



# PyAudio



# PyAudio - Installation

## Brew


```
brew install portaudio
```

## PyPI

```
pip install pyaudio
```

## Conda

```
conda install -c conda-forge pyaudio
```



# Recording & Visualizing Frequencies

Goals:

Capture sound by a microphone in real-time

Visualize the frequencies present in real-time



# Recording & Visualizing Frequencies

1. Capture the sound in fixed-size segments (200 ms)
2. For each segment, plot the frequency distribution in real-time
  1. Compute magnitude  $X$  of FFT of the segment and frequency values (in Hz) in a separate array `freqs`
  2. Downsample  $X$  and `freqs`, so that we keep very few frequency coefficients to visualize
  3. Compute total segment's energy (to normalize against the maximum width of the frequency visualization)
  4. Bar Plot downsampled frequency energies  $X$  for all frequencies

# Recording & Visualizing Frequencies

Capture the sound in fixed-size segments (200 ms)

```
pa = pyaudio.PyAudio()  
stream = pa.open(format=pyaudio.paInt16, channels=1, rate=fs,  
                  input=True, frames_per_buffer=int(fs * buff_size))
```

# Recording & Visualizing Frequencies

Compute magnitude  $X$  of FFT of the segment and  
frequency values (in Hz) in a separate array `freqs`

```
X = np.abs(scipy.fft(x))[0:int(seg_len/2)]  
freqs = (np.arange(0, 1 + 1.0/len(X), 1.0 / len(X)) * fs / 2)
```

# Recording & Visualizing Frequencies

Downsample  $X$  and  $\text{freqs}$ , so that we keep very few frequency coefficients to visualize

```
wanted_step = (int(freqs.shape[0] / wanted_num_of_bins))  
freqs2 = freqs[0::wanted_step].astype('int')  
X2 = np.mean(X.reshape(-1, wanted_step), axis=1)
```

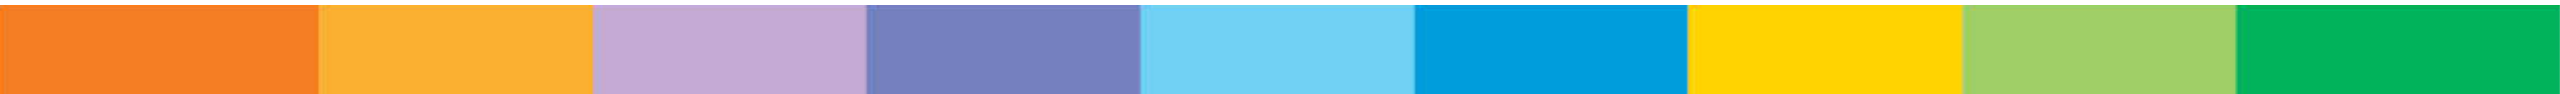
# Recording & Visualizing Frequencies

Bar Plot downsampled frequency energies  $X$  for all frequencies

```
fig = tpl.figure()
fig.barh(X2, labels=[str(int(f)) + " Hz" for f in freqs2[0:-1]],
         show_vals=False, max_width=max_width_from_energy)
fig.show()
```

# Recording & Visualizing Frequencies

Demo





# Librosa

<https://librosa.org/doc/latest/index.html>

# Librosa - Installation

## PyPI

```
pip install librosa
```

## Conda

```
conda install -c conda-forge librosa
```

## Source

```
tar xzf librosa-VERSION.tar.gz
```

```
cd librosa-VERSION/
```

```
python setup.py install
```

# Librosa - Submodules

**beat** - Estimating tempo and detecting beat events

**Core** - Load audio, compute spectrograms and other analysis

**Decompose** - Harmonic-percussive source separation (HPSS)

**Display** - display and visualization routines

**Effects** - Pitch Shifting, Time-stretching

**Feature** - Chromagrams, Mel Spectrogram, MFCC

**Filters** - Filter-bank generation (Chroma, pseudo-CQT, CQT)

**Onset** - Onset detection and Onset strength computation

# Librosa - Quick Example (Beats)

Importing library

```
import librosa
```

Get the audio file

```
filename = librosa.example('nutcracker')
```

Load the audio file

```
y, sr = librosa.load(filename)
```

Beat tracker

```
tempo, beat_frames = librosa.beat.beat_track(y=y, sr=sr)
```

Tempo

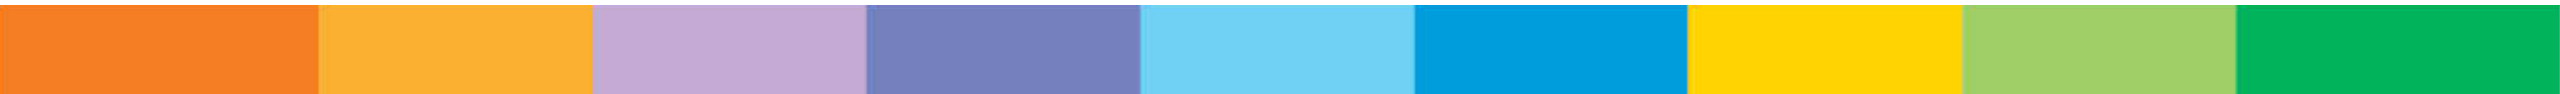
```
print('Estimated tempo: {:.2f} beats per minute'.format(tempo))
```

Frame numbers into  
timestamps

```
beat_times = librosa.frames_to_time(beat_frames, sr=sr)
```



# Beats Generator



# Beats Generator - Tempo tracking


Task of automatically estimating a song's temp (in beats per minute) directly from the signal using librosa library

Input: Mono audio file

Output: Stereo file

Left channel: original song

Right channel: "beep" sound mimicing the tempo of the song



# Beats Generator - Tempo tracking

```
[Fs, s] = wavfile.read('data/music_44100.wav')
tempo, beats = librosa.beat.beat_track(y=s.astype('float'), sr=Fs, units="time")
beats -= 0.05
s = s.reshape(-1, 1)
s = np.array(np.concatenate((s, np.zeros(s.shape)), axis=1))
for ib, b in enumerate(beats):
    t = np.arange(0, 0.2, 1.0 / Fs)
    amp_mod = 0.2 / (np.sqrt(t)+0.2) - 0.2
    amp_mod[amp_mod < 0] = 0
    x = s.max() * np.cos(2 * np.pi * t * 220) * amp_mod
    s[int(Fs * b):
      int(Fs * b) + int(x.shape[0]), 1] = x.astype('int16')
wavfile.write("data/music_44100_with_tempo.wav", Fs, np.int16(s))
```

# Beats Generator

Sample 1



Sample 2





# PyAudioAnalysis

<https://github.com/tyiannak/pyAudioAnalysis>

# PyAudioAnalysis - Installation

## PyPI

```
pip install PyAudioAnalysis
```

## Conda

```
conda install -c conda-forge
```


```
PyAudioAnalysis
```

## Source

```
git clone https://github.com/tyiannak/pyAudioAnalysis.git
```

```
cd pyAudioAnalysis
```

```
pip install -r ./requirements.txt
```



# Further Reading

Basic Concepts in Audio Processing

<https://www.pythonforengineers.com/audio-and-digital-signal-processingdsp-in-python/>

<https://www.ee.iitb.ac.in/student/~daplab/publications/chapter9-prao.pdf>

Librosa

<https://librosa.org/doc/latest/index.html>

PyAudioAnalysis

<https://github.com/tyiannak/pyAudioAnalysis>

Pydub

<https://github.com/jiaaro/pydub>